

TOPIC 11 ASSESSMENT REPORT

Security, Rate Limits & Production Scaling — Retell AI

Prepared by: Shaik Mohammad Shaheed **Agent:** Trainee_Sarah_Receptionist

Evidence update: 6 October 2026

1. Objective

Document the security and scaling settings reviewed in Retell AI and record the observed result of one Playground test call for PII handling. This report distinguishes settings visible in the dashboard from behavior actually verified during the call.

2. Configuration Reviewed

Area	Observed configuration / status
Data storage	"Everything except PII" selected; retention configured to 30 days.
PII redaction	14 redaction categories configured.
Keypad / DTMF	User Keypad Input Detection enabled; timeout 2.5 seconds; digit limit enabled and set to 4.
API keys	A masked Secret Key named "Webhook" was visible. No key rotation or revocation performed.
Limits	Concurrent Calls Limit displayed as 20. Dashboard also showed LLM token and telephony CPS limits.
Reliability	Stable Server feature remained disabled; paid add-on was not enabled.
Webhook	A webhook endpoint was configured. Webhook test was not run.

3. Playground PII Test — Observed Result

One Retell Playground web call was completed on 6 October 2026. Call History showed a duration of 0:29 and a displayed cost of \$0.090 (the audio player showed 0:28). The transcript displayed placeholders such as [person name 1] and [credit card 1] instead of the entered name and credit-card digits. The agent warned the caller not to share credit-card details and redirected them to appointment or billing assistance.

Observed: Transcript-level masking is visible for this test utterance. **Not yet verified:** whether the corresponding audio segment was muted. The screenshot alone cannot establish audio muting.

4. Evidence and Deliverables

Existing repository evidence documents data storage, PII categories, webhook settings, keypad/DTMF settings, masked API keys, limits, and reliability settings. The new Call History screenshot provides additional evidence of duration/cost, the masked transcript, and the agent's warning. Audio behavior should only be reported after separately checking the recording segment.

5. Tests Not Performed

No API key rotation/revocation test, 110-request rate-limit test, 21-call concurrency test, webhook delivery test, retention deletion test, or encryption-in-transit/at-rest audit was performed. These remain pending; configuration screenshots alone do not prove runtime behavior. No claim is made that audio muting, HTTP 429 behavior, encryption algorithms, or load handling have been independently verified.

6. Recording and Repository

Loom demonstration: [Open Loom video](#)

GitHub repository: [Open GitHub repository](#)

7. Final Status

The configuration review and one Playground PII transcript check are documented. Transcript masking is visible in the evidence; audio muting and the listed load/security integration tests remain unverified. This is an evidence-based assessment report, not a certification of production security.

